

ConstellationCV: Obtaining 3D Information from 2D Images Efficiently using Feature Extraction and Laser Dot Projection

Pratham Gandhi (pratham_gandhi@horacemann.org),
Samuel Schuur (samuel_schuur@horacemann.org)

Bronx, New York

Abstract

In this paper, the development and application of a system to computationally inexpensively determine a three dimensional model of a space from only a single image is presented, using applied machine learning and simple hardware components. Using feature classification of the relevant objects in the original scene, in collection with a physical grid of laser dots refracted across the scene, the system is able to quickly judge the nature of the object, its position in relation to other objects around it, and distance of those objects from the camera, traits which are commonly compromised on comparable approaches. The system is x times more efficient on average than current industry leading approaches, and has potential applications in various fields requiring rapidly updating spatial awareness and environment generation, including autonomous vehicles, defense, and architecture.

1. Introduction

The current most popular approach to creating a three dimensional informational estimate from a two dimensional image is triangulation, or binocular, stereo vision approach. This approach take several images of the same scene from different locations and angles in order to generate a spatial awareness. Computer stereo vision, however, is difficult to implement in practical applications in our increasingly evolving digital world, due to the limiting factor that they require multiple images, the acquisition of which is sometimes impractical. Additionally, they can sometimes be quite slow, in that they need to first identify common features in the different photos, find out the positional relation of those features, and then only begin to start seeing the greater picture of the full space. The purpose of Constellation was to solve both of these problems and create a flexible system which can adapt and continue to be functional in many environments.

We chose to approach this problem first through the somewhat traditional technique of using artificial neural networks, mathematical functions modeled after nature which specialize in taking in large amounts of input to produce outputs, to identify features in a two dimensional image. Neural networks are perfect for navigating the fuzzy problems of feature extraction from images. Where our system differs is in its use of a grid of refracted laser points into the scene to judge the distances and orientations of various objects in relation to the camera. This allows superior environment generation in a shorter amount of time, due to the omission of the multiple-image stereo aspect. This approach and its implemented system was named Constellation for the star-like appearance of the laser dot grid it relies on.

2. Background

2.1. Binocular Computer Stereo Vision

In order to perceive depth the human eye uses many visual cues, both monocular and binocular. Of these visual cues, stereopsis, the perception of depth resulting from the brain comparing the differences in images produced by both eyes, is in many ways one of the easiest depth perceptions methods to mimic through software, as it is the least reliant on external sources of stimuli and human experience. This is what Stereo Computer Vision attempts to do; building depth maps by comparing scenes from multiple vantage points. More specifically Binocular Computer Stereo Vision does this by comparing

the images from two cameras and by extracting and comparing common features from both images it builds a depth map of the scene in question.

2.1.1. Removing Distortion from Images

2.2. Structure-From-Motion Pipelines




Figure 1: SFM Camera Setup Diagram

Structure-from-motion (SFM) pipelines operate on a simple premise; finding common anchor points among several images of the same subject, and using triangulation to estimate where those points lie in relation to the camera.

2.2.1. Shortcomings

One of the major issues in using SFM pipelines in practical applications where rapidly updating three dimensional estimations are required is that the models produced by SFM pipelines are lacking a scale factor, that is, there is no sense accurate of scale or distance in the model. This is one of the issues Constellation aims to address.

2.3. Neural Networks

An Artificial Neural Network (ANN) is a computational model inspired by biological networks in the human brain which process large amounts of information.

2.3.1. A Single Neuron

The basic unit of a neural network is a neuron. A neuron is essentially a node in the graph which represents the neural network. Each neuron in an ANN, similar to a biological neuron, receives input from other neurons. Each input has an associated weight W , which is adjusted based on the importance of its corresponding input in the large picture of all the inputs to a neuron. Once it has acquired its inputs, a neuron will apply an activation function f , as shown below, to the weighted sum of its inputs to produce an output.




Figure 2: Artificial Neuron Diagram

$$\text{input to activation} = \left(\sum_{i=1}^n (x_i \times w_i) \right) + b$$

$$\text{activation } f(x) = \frac{1}{1 + e^{-x}}$$

$$\text{output} = \frac{1}{1 + e^{-\left(\left(\sum_{i=1}^n (x_i \times w_i) \right) + b \right)}}$$

The function f is a non-linear activation function. The role of the activation function is to create non-linearity in the output of a neuron, which is key because most real-world data is non-linear and the purpose of neurons is to learn how to represent that data. Every activation function takes a single input and performs a mathematical operation on it in order to produce the final output. Common activation functions include:

- Sigmoid: takes a value and squashes it to range 0-1. $\sigma(x) = \frac{1}{1+e^{-x}}$

- tahn: takes a value and squashes it to range $[-1,1]$. $\text{tahn}(x) = 2\sigma(2x) - 1$
- ReLU: stands for Rectified Linear Unit. Takes a value and replaces all negative values with 0. $f(x) = \max(0, x)$

Constellation chooses to use sigmoid as the primary activation function due to its gradient nature, fixed range, and the fact that it is easier to differentiate values closer to one asymptote or the other.

2.3.2. Feedforward Networks

The feedforward neural network is the simplest and most widely applied type of ANN devised. Feedforward nets contain several neurons sorted into layers. Nodes from adjacent layers have connections, represented as edges on the graph, between them. Every connection, as discussed earlier, has a weight associated with it.




Figure 3: An example of a feedforward neural network

(a) Note that an empty circle on an edge denotes a source and a filled circle on an edge denotes a sink.

A feedforward network can consist of three types of nodes:

1. Input Nodes - Provide information from the outside world to the network, the set of which is called the "Input Layer". None of these nodes perform any computation.
2. Hidden Nodes - Have no direct connection with outside world and only perform computations and transfer information from the input nodes to the output nodes. A set of these nodes is called a "Hidden Layer". Feedforward nets can have any number of hidden layers.
3. Output Nodes - Collectively called the "Output Layer", these nodes perform computations and also transfer information from the network to the outside world.

In a feedforward network, information only moves forward, through the input nodes, hidden nodes, and then finally on to output nodes. There is no cycling or looping back through the network. Below is an abridged psuedocode version of the feedforward algorithm implemented in Constellation:

Algorithm 1 Feed-Forward Algorithm

```

1: procedure FEEDFORWARD:
2:   create an empty list of all layers' output
3:   loop for each layer in network:
4:     create empty list of this layer's outputs
5:     biased input  $\leftarrow \sum_{i=1}^n (x_i \times w_i) + b$ 
6:     loop for each neuron in layer :
7:       neuron output  $\leftarrow \sigma(\text{biased input})$ 
8:       layer's outputs  $\leftarrow$  layer's outputs + neuron output
9:     all outputs  $\leftarrow$  all outputs + layer's output
10:  return all outputs

```

2.3.3. Training Using Back-Propagation

It is possible to train a Multi-Layer Perceptron (MLP) Neural Network, a network which has at least one hidden layer, using a technique called back-propagation. Back-Propagation is a type of training approach which is called supervised learning, which means that the network is learning from labeled training data, where there is data is fed in, accompanied by classifications for each set of data which the network is told are correct. Each connection between nodes has a weight, which, as mentioned earlier, dictates the relevance of a specific input in the greater scope of all the inputs to a given neuron. The goal of learning is to assign correct weights to these edges in order to provide the most accurate classifications.

All weights are randomly assigned to begin with. Then, one by one, every training set is passed through the neural network, and, depending on how vast the difference desired and actual output of each layer, the error is propagated back to the previous layer. The error is noted and each weight is adjusted accordingly. This process is repeated using the provided data set until the network is outputting acceptably close classifications to the labeled

classifications. Once this level is reached, the network can be passed previously unseen data and can use its trained weights to produce a classification.

Below is an abridged psuedocode version of the backpropagation algorithm implemented in Constellation:

Algorithm 2 Backpropagation Algorithm

```

1: procedure BACKPROPAGATE(NETWORK, INPUT VECTOR, TARGETS):
2:   hidden outputs, outputs  $\leftarrow$  feed_forward(network, input vector)
3:   loop for each output:
4:     delta output  $\leftarrow \dot{\sigma}(\text{output}) * (\text{output} - \text{target})$ 
5:   loop for each output:
6:     adjust weights of the  $i$ th neuron
7:   loop for each hidden output: // hOutput = hidden output
8:     hidden delta  $\leftarrow \dot{\sigma}(\text{hOutput}) * (\text{hOutput} - \text{next adjusted input})$ 
9:   loop for each hidden output:
10:    adjust weights of the  $i$ th hidden neuron

```

3. Approach

3.1. Determining a 3D Model of an Object From a Video

Of course, creating a three dimensional model of a space is not possible without knowing three dimensional information about the individual objects which are to be extracted from the two dimensional image and placed in the three dimensional scene. Constellation takes a slightly unconventional approach to this problem in that it implements the traditional approach it is trying to improve, a structure from motion pipeline. However, the goal of producing a long-run rapidly updating computationally inexpensive system is still being satisfied. While they are not practical in situations with large scenes containing many objects being generated relatively fast, SFMs perform at acceptable levels for smaller datasets, as their complexity is exponential. Constellation only employs the comparatively slow SFMs at non time-critical moments; when training the network to recognize an object. While training itself to recognize an object in that image, Constellation concurrently uses a structure from motion pipeline to construct a three dimensional model of that object and stores that model.

As explained in the background (2.2), structure from motion pipelines use a series of images of the same object from different perspectives to create the three dimensional model of something. The pipeline can either be presented with the locations of the camera when the different pictures were taken, or it can use some computations to figure that for itself. Once it has established the locations from which each of the several images were taken, the pipeline extracts common anchor points by finding distinct and easily distinguishable features in every image, and then matching those anchor points from one view of the object to another in order to determine the location of that point in a three dimensional space, taking into account the difference in position within the images of those anchor points. Evidently, an SFM pipeline requires several images of an object in order to establish a 3D representation of the object, and Constellation achieves this by dissecting a video clip taken which captured as many unique views of the object as possible into a series of images, precisely 3 images from each second of video. Not only will these images be useful in constructing that model of this object, they serve a second purpose as they are also used to train Constellation to detect the object in images.

Due to the sparing and time-sensitive use of SFMs described earlier, Constellation has already learned what the model of individual objects in a scene look like at the time when it has to construct a model of a full scene, and thus saves precious compute time while it is generating models of the full scene from a 2D image on the fly.

3.2. Feature Detection and Classification

The central problem of detecting various features within an image is one of classification. There exist several different approaches or algorithms for classification, including decision trees, naive, bayes, and KNN. However, only two classification models stood out for our particular application of object recognition; support vector machines (SVMs), and neural networks. SVMs operate by determining an optimal classification line, which separates training data of two different types into two distinct sections, and performs classification by determining which side of that line a set of input data falls. While SVMs have a higher accuracy in general, and are more tolerant to redundant and irrelevant attributes, they require on average three times more training samples to accurately classify features than neural networks, and still perform evenly with neural networks with regards to the speed of classification, speed of learning, and tolerance to highly interdependent attributes. This system's

applications require rapidly updating environments, which thus require rapid computations and rapid training. These reasons make SVMs impractical for use in Constellation, which means the best choice for Constellation’s object detection and classification approach was neural networks, which could easily be pretrained and then adjusted during operation. With neural networks identified as the best approach to classifying various features in an image, the problem of how to actually implement and use a neural network to perform this task presents itself. The approach to detecting objects in an image has three major steps; training, iteration, and classification.

As discussed earlier (3.1), during the construction of models of individual objects, a series of images are generated from videos of each of the objects. These same images can be used to train the neural network using the backpropagation algorithm presented in the background (2.3.3). Each of these training images are greyscaled and then broken down into an array of values between 0 and 1, 0 representing pixels which are completely dark, and 1 representing pixels which are completely white. This array is then fed forward through the neural network and then that output is backpropagated, training the network more with every image fed through. Once the network is trained, it can be used to find those features in an image. In reality, the second two steps of the object detection process, iteration and classification, mesh into one more general process. Given the height and width of the objects the network was trained on, Constellation iterates over the image in question and employs a "cookie cutter" strategy of cutting out smaller images of the size of the training images at regular intervals within the image, and checks each of these smaller images for the presence of one of the objects the network was trained on. If the network’s confidence that one of these images is present exceeds a set threshold value, the network classifies this "sub-image" as the location of the classified object.

The same approach can be used to find the laser dots which, as discussed in the following sections, are quite vital to the distance estimation needed to construct the model of the scene. Several images of these dots were taken in various environments and a separate network was trained on these images, using a similar approach as that of object classification. The training of the network used to detect the dots differs in that, instead of greyscaling the images as done in the other approach, colored images are considered, and instead of arrays of values between 0 and 1, the network is provided an array of arrays of red, green, blue (RGB) values. The classification of all the laser dots in an image is run then in the same manner as that of objects, and

is run concurrently with that process, as both processes employ separately trained networks and other resources.

3.3. Distance Estimation

Constellation is unique from other systems in its approach of judging the distance of various objects within its captured scene from the camera. Most other approaches might use computationally intensive math and thus have very slow running times. Constellation, instead, uses efficient neural network driven image classification, in combination with simple statistical analysis to perform the same task. The underlying principle which allows this approach to function is that, when refracted through a particularly shaped crystal, the light from a single laser is scattered in such a manner that, as it moves further in distance from its source, it strays further in one direction or the other from its origin. Thus, when it is intercepted by another object at a certain distance, the amount the light at that distance has moved away from the position it would be at if it was intercepted directly in front of the source should be somewhat representative of the numerical distance travelled by that light, and additionally, the distance between individual points of refracted light should also be indicative of the distance from the camera, as each refracted beam of light would stray further from its neighbors.

Constellation’s object detection and classification subsystem provides the image coordinates of the center of the object it detects. This is the foundation point used for judging the distance within the image between the laser dots refracted across the figure. In order to find the distance of the object from the camera, the distance between the dots projected onto the object has to be found. Constellation does this by obtaining a sorted list of all the points in the image, a process which in the worst case is $O(n \log n)$, due to the choice to employ a modified Timsort algorithm, but due to the nature of Constellation’s object detection implementation, trends nearer to $O(n)$. Following processing and sorting the list of all the dots found in the image, Constellation proceeds to conduct a modified binary search through the list in an effort to find the laser dot closest to the center of the object, taking $O(\log n)$ in both the worst and average cases. Once that has been acquired, the system takes an iterative approach to finding as many possible non-distorted dots on the surface of the object, calculating the euclidean distance between the dot closest to the center of the object and the 4 nearest dots, and checking if they are within the shape of the object determined during the process of training Constellation to recognize the object. If all 4 of the closest

dots are within the shape of the object, Constellation proceeds to check the next 4 closest dots, and so on. The entirety of these processes assembles an array of all the representative dots on the visible surface of the object, and that array is used to find the average euclidean distance between each all the points on the surface of the object.

As mentioned earlier, the distance between the refracted dots projected onto an object is theoretically representative of the distance of that object from the camera which captured the image. Several approaches were considered when determining the most efficient and accurate method to extract information about an object’s distance from the camera, including finding an as-accurate-as-possible multiple-degree polynomial modeling of the distance between dots vs. the distance from the camera, linear modeling, training a neural network to predict the distance, and a k-nearest neighbors (KNN) approach. Firstly, we considered multiple degree polynomials to mathematically model the correspondence between average distance between dots on the surface of an object and the distance of that object, however, polynomials over a degree of 2 often result in over-fitting, or creating a model which too closely or too accurately represents the training set. This might seem like an advantage of this approach, however, over-fitted models never generalize well to datasets beyond the training set, and thus, as shown in figure 5, actually lose accuracy. So, the rational next step after multiple degree polynomials is linear modeling. A linear function representing the relationship would probably be fairly accurate and abstract well enough to general datasets, however, linear functions exhibit another phenomenon called "underfitting", which means that the model is neither very accurate with the training set nor with general sets. Having ruled out functional modeling all together, the focus was shifted towards models grounded in statistical science and machine learning. Having already heavily implemented them in other parts of Constellation, neural networks were considered for the tasks, but it was determined that they were unnecessary due to their relatively extensive training data necessity and comparatively slow classification speed. Another machine learned based model, SVMs, were considered, but SVMs can only really produce one decision boundary in a dataset well, and therefore cannot be abstracted to classify more than 2 possibilities or labels easily (Weston, Jason, and Chris Watkins, 1998). All of these reasonings led to the choice of a KNN model, a model which finds the k nearest data points to the one which is to be classified, and essentially lets them "vote" on the classification of the data point based on their own classifications, therefore choosing

the majority classification of the points near the target point. The graph in figure 5 graphically illustrates the advantages and disadvantages of each of the viable considered models in this particular task on a sample dataset.

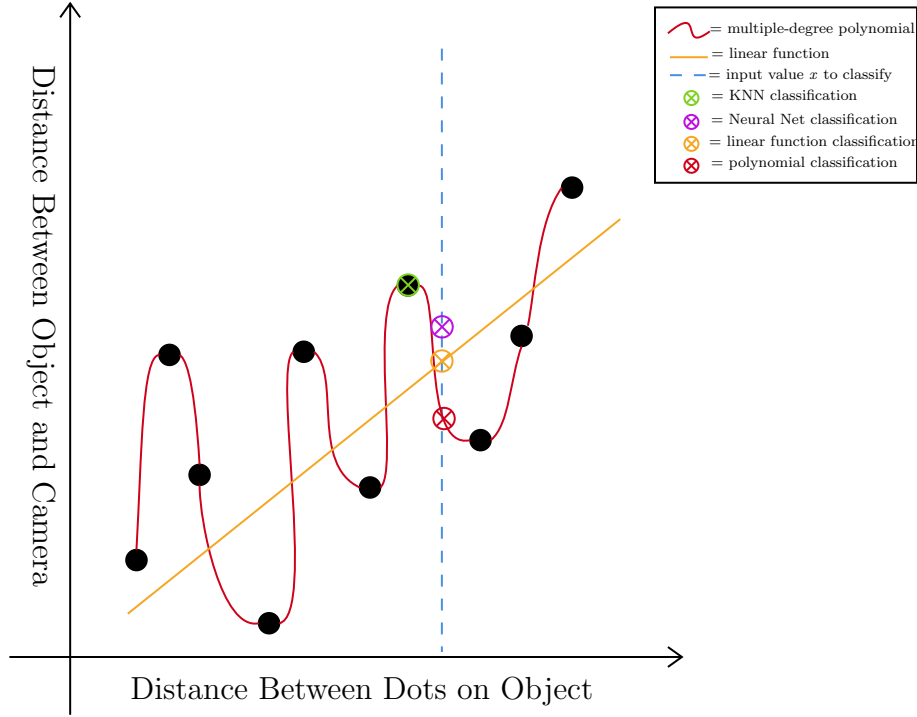


Figure 5: Different Classification Approaches on Sample Set

A KNN model's accuracy grows near-exponentially with the addition of more labeled training data, as its training data is essentially the entire model. Constellation currently operates on limited range, due in part to the limited visibility of the laser dots as they move farther from the source, but also due to the limitation of available training data. As of now, a relatively large training set with over x data points has been gathered, and is effectively used to predict distance within a range of x . As described earlier, the model predicts the distance of an object depending on the spread of the laser dots on the object's surface by finding the values of the 6 nearest data points and letting them vote on an appropriate distance for the object. Assume there is a defined function `majority_vote(labels)` which returns the majority vote of k labels and, in the case of ties, recursively reduces k until there are no more ties. Given that, the following is an abridged pseudocode segment of

Constellation’s implementation of a KNN model illustrating the logic behind the model:

Algorithm 3 KNN Classification Algorithm

```

1: procedure KNN_CLASSIFY(K, LABELED POINTS, NEW POINT):
2:   by distance  $\leftarrow$  sorted labeled points by distance from new point
3:   create empty list of k nearest
4:   loop for first k points in by distance:
5:     k nearest  $\leftarrow$  k nearest + point
6:   classification  $\leftarrow$  majority_vote(k nearest)
7:   return classification

```

4. Implementation

4.1. Hardware

4.1.1. Refracting Grid of Laser Points and Camera Rig

4.1.2. Computers

While most practical industry or even hobby applications of the concepts presented in Constellation and the library itself will take place on far more powerful machines than those Constellation was developed and tested on, in its testing, Constellation still performed well compared to other approaches to the same problem. For compatibility assurance, Constellation was tested primarily on three separate machines, each with drastically different components, and tested with two operating systems.

The first testing machine ran a copy of Apple’s macintoshOS operating system. The machine had 2.3 Gigahertz Intel Core i5 processor, which allowed the processor to run at a clock rate of 2.3 billion cycles per second. Additionally, the machine had 8 Gigabytes of 2133 Megahertz LPDDR3 memory, and a 1536 megabyte Intel Iris Plus Graphics 640 graphics card. The second testing machine was operating on the Microsoft Windows operating system.

As shown below, the system was still able to perform well on the Apple computer, despite comparatively inferior components, due to a concurrency modification that is still in beta, which allowed the computer to distribute the computations of certain processes across two processors, thereby decreasing the load of each individual processor.

4.2. Software

4.2.1. Determining a 3D Model of an Object From a Video

4.2.2. Object Oriented vs. Procedural Neural Network Design

4.2.3. Feature Extraction

In constellation's implementation, there are three main parts of extracting features from a two dimensional image; the first is training our system to recognize the object




Figure 6: Number of Matches Found and Speed of Matching (seconds) vs. Confidence Level For Varying Amounts of Training Images

4.2.4. Distance Estimation

4.2.5. Environment Model Generation

4.2.6. User-Facing Application

5. Results

5.1. Algorithmic Analysis

Constellation's increased speed is due in part to its superior algorithmic efficiency. First consider its core object detection system, a neural network.

The network is pre-trained, so in the analysis of its efficiency, the training steps, including the backpropagation algorithm used to fine tune its weights, will be ignored, as at run-time, they will not play any role in the time and number of calculations it takes to determine the nature of an object. Thus, the complexity of identifying an object is only the complexity of feeding the data from the image through the neural network and producing an output. The derivations and declarations of variables used to figure the complexity of the feed-forward algorithm are:

$$\# \text{ layers} = \text{constant } c \text{ (in our implementation } c = 5)$$

$$\frac{\# \text{ neurons}}{\text{layer}} = \sqrt{n} \text{ for } n \text{ elements in an input feature vector}$$

$$\frac{\# \text{ calculations}}{\text{neuron}} = n + 5$$

$$\# \text{ times object detection called} = (w - \sqrt{n})(h - \sqrt{n})$$

(for image width w , and image height h ;

in our implementation $w = 1080$, $h = 720$)

Thus, it can be determined that the number of calculations required to determine the classification of a set of n inputs is $5\sqrt{nn} + 25\sqrt{n}$, and further, the complexity of classifying all objects in an image is $O(n^2)$.

To perform similar tasks, a fairly naive implementation of a Structure-From-Motion pipeline would require $15n^3 + c$ calculations, and a stereo vision system implementation would require *find this* calculations, both far more than Constellation's mere *find this* needed calculations for fairly standard object image size of $n = 100$ pixels.

5.2. Implemented Execution Time

5.3. Environment Accuracy

6. Conclusion

6.1. Advantages Over Similar Systems

6.2. Shortcomings and Future Improvements for Constellation

- need to pre-train the network on objects before Constellation can put it in the environment accurately

- rez extract key anchor points such as pre-trained corner or vertices, and interpolate a reasonably applicable mesh over those vertices to form objects so that you don't have to train full and exact objects
- non-visual dot projection (infrared or something else)
- increase time efficiency by concurrently carrying out object detection and dot detection operations (weight the pros and cons of distributed computations at that scale and their intensiveness on cpu, also weave in gpu acceleration)
- range, increase by stronger lasers and more training data for KNN model

Appendix A. References

- [1] M. Xu, X. Liu, Y. Liu, F. X. Lin, Accelerating convolutional neural networks for continuous mobile vision via cache reuse, arXiv preprint arXiv:1712.01670 (2017).
- [2] M. Pollefeys, L. Van Gool, M. Vergauwen, F. Verbiest, K. Cornelis, J. Tops, R. Koch, Visual modeling with a hand-held camera, International Journal of Computer Vision 59 (2004) 207–232.
- [3] J. Li, M.-T. Luong, D. Jurafsky, E. Hovy, When are tree structures necessary for deep learning of representations?, arXiv preprint arXiv:1503.00185 (2015).
- [4] I. Tsochantaridis, T. Joachims, T. Hofmann, Y. Altun, Large margin methods for structured and interdependent output variables, Journal of machine learning research 6 (2005) 1453–1484.
- [5] J. Weston, C. Watkins, Multi-class support vector machines, Technical Report, Citeseer, 1998.
- [6] K. Crammer, Y. Singer, On the algorithmic implementation of multiclass kernel-based vector machines, Journal of machine learning research 2 (2001) 265–292.
- [7] M. Potmesil, Generating octree models of 3d objects from their silhouettes in a sequence of images, Computer Vision, Graphics, and Image Processing 40 (1987) 1–29.

Appendix B. Open Source Code and Library

The ConstellationCV project was created and built as a research project to explore new approaches to an old problem. As such, most of the code and documents produced in the research project have been open-sourced and are accessible by anyone. All of the production code has been formatted into an open library structure to allow for easy use of it, as well as to encourage the modification and adaptation of the principles for specific uses. The authors of this paper welcome any contributions to or thoughts about the project, and would additionally encourage those interested to fork the repository on GitHub and attempt improve it in any way possible. All of the code can be found at:

`https://github.com/ConstellationCV/Constellation`

Additionally, for those interested, all documents produced in the process of developing Constellation, including research logs and all the source files for this paper, have been open sourced and can be found at:

`https://github.com/ConstellationCV/Documents`